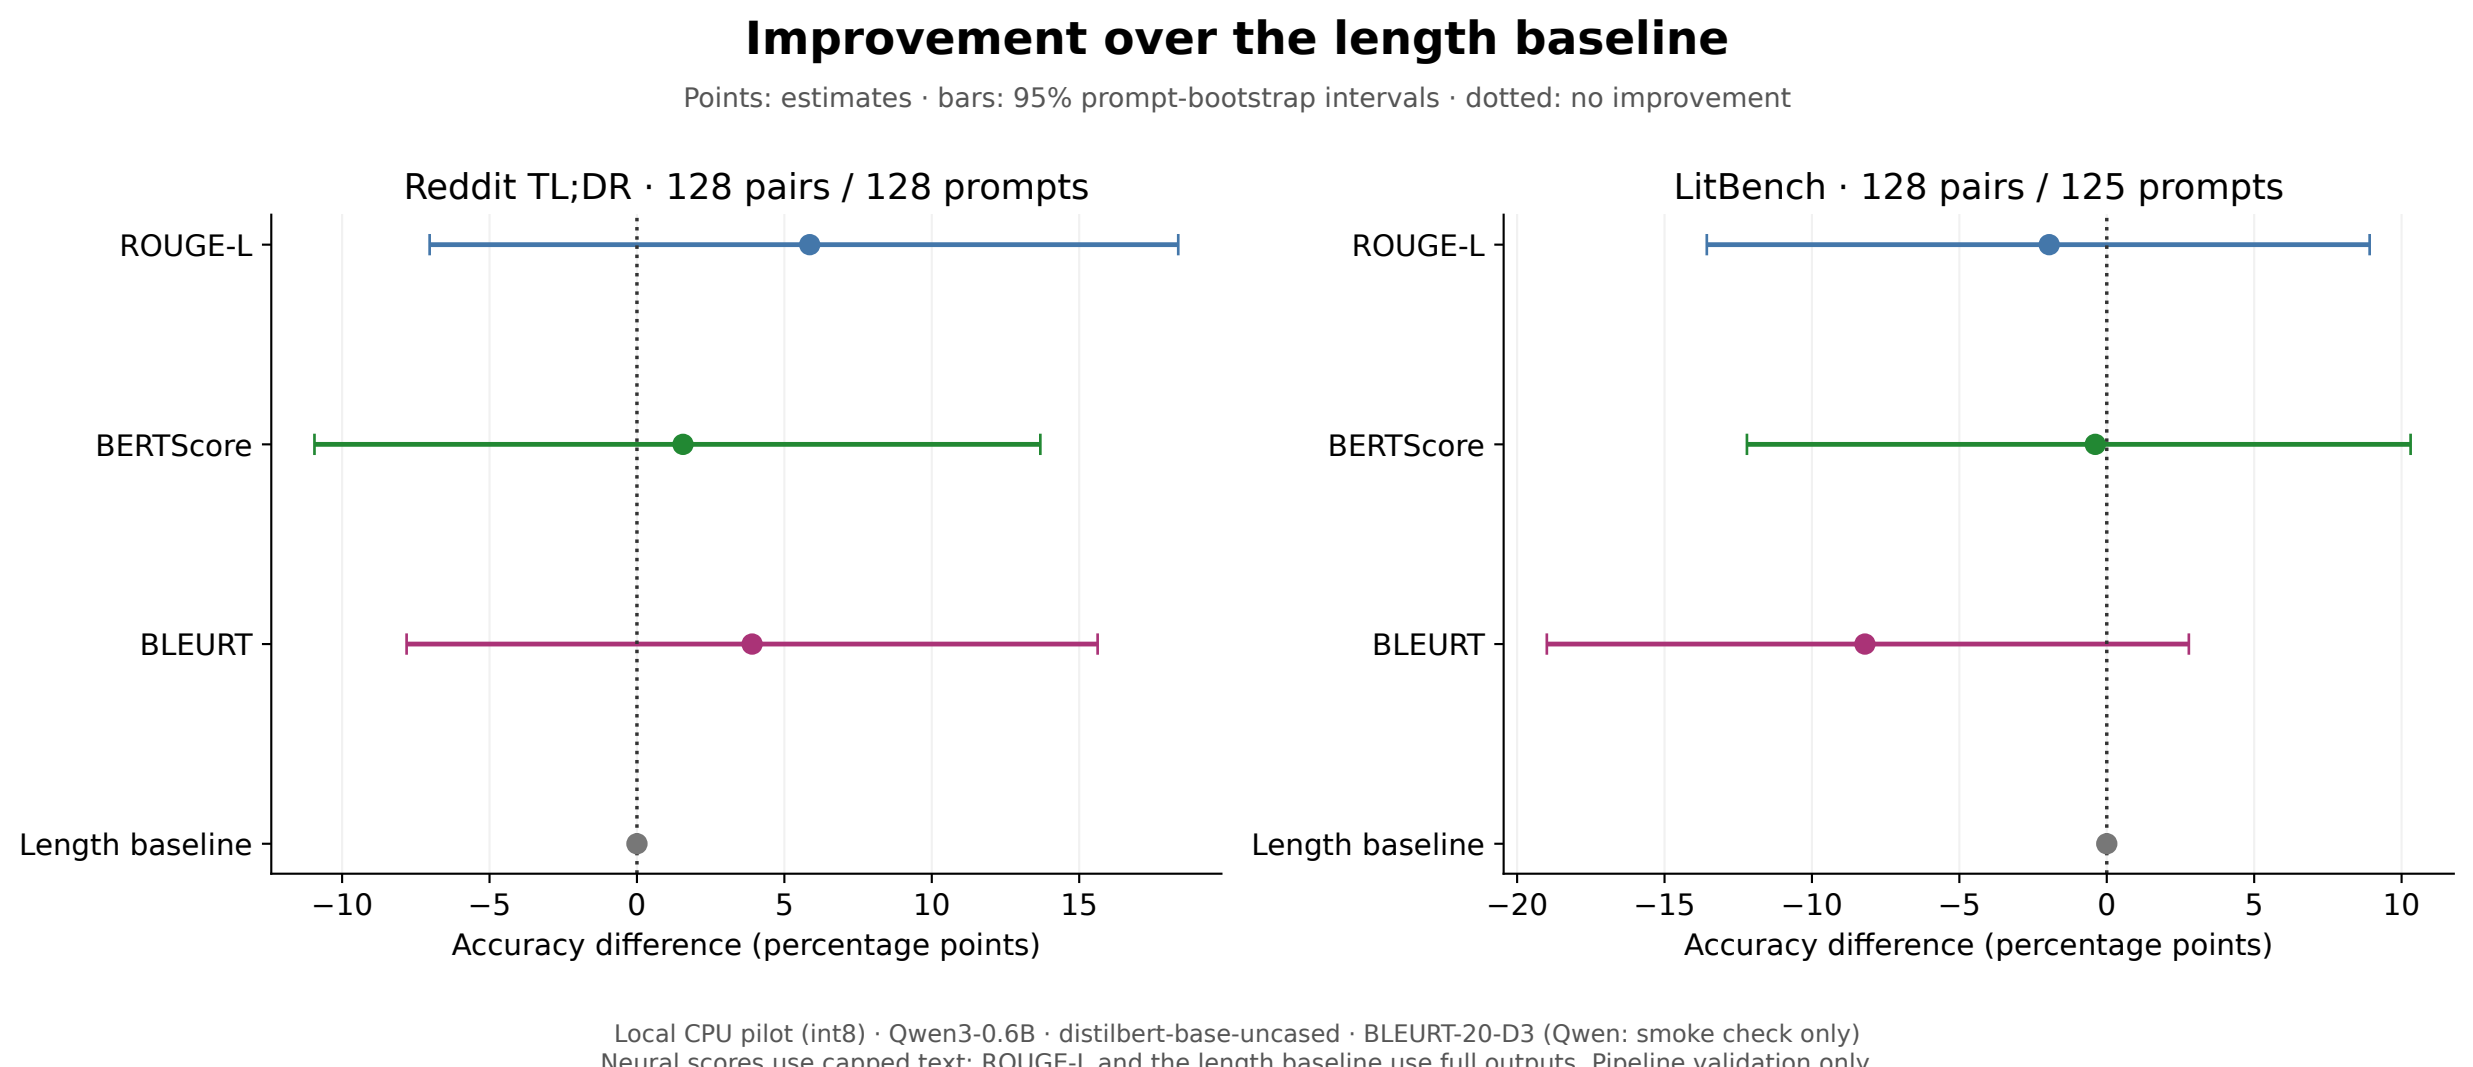
Improvement over the length baseline
Points: estimates · bars: 95% prompt-bootstrap intervals · dotted: no improvement
Reddit TL;DR · 128 pairs / 128 prompts
LitBench · 128 pairs / 125 prompts
ROUGE-L
BERTScore
BLEURT
Length baseline
−10
−5
0
5
10
15
−20
−15
−10
−5
0
5
10
Accuracy difference (percentage points)
Accuracy difference (percentage points)
Local CPU pilot (int8) · Qwen3-0.6B · distilbert-base-uncased · BLEURT-20-D3 (Qwen: smoke check only)
Neural scores use capped text; ROUGE-L and the length baseline use full outputs. Pipeline validation only.