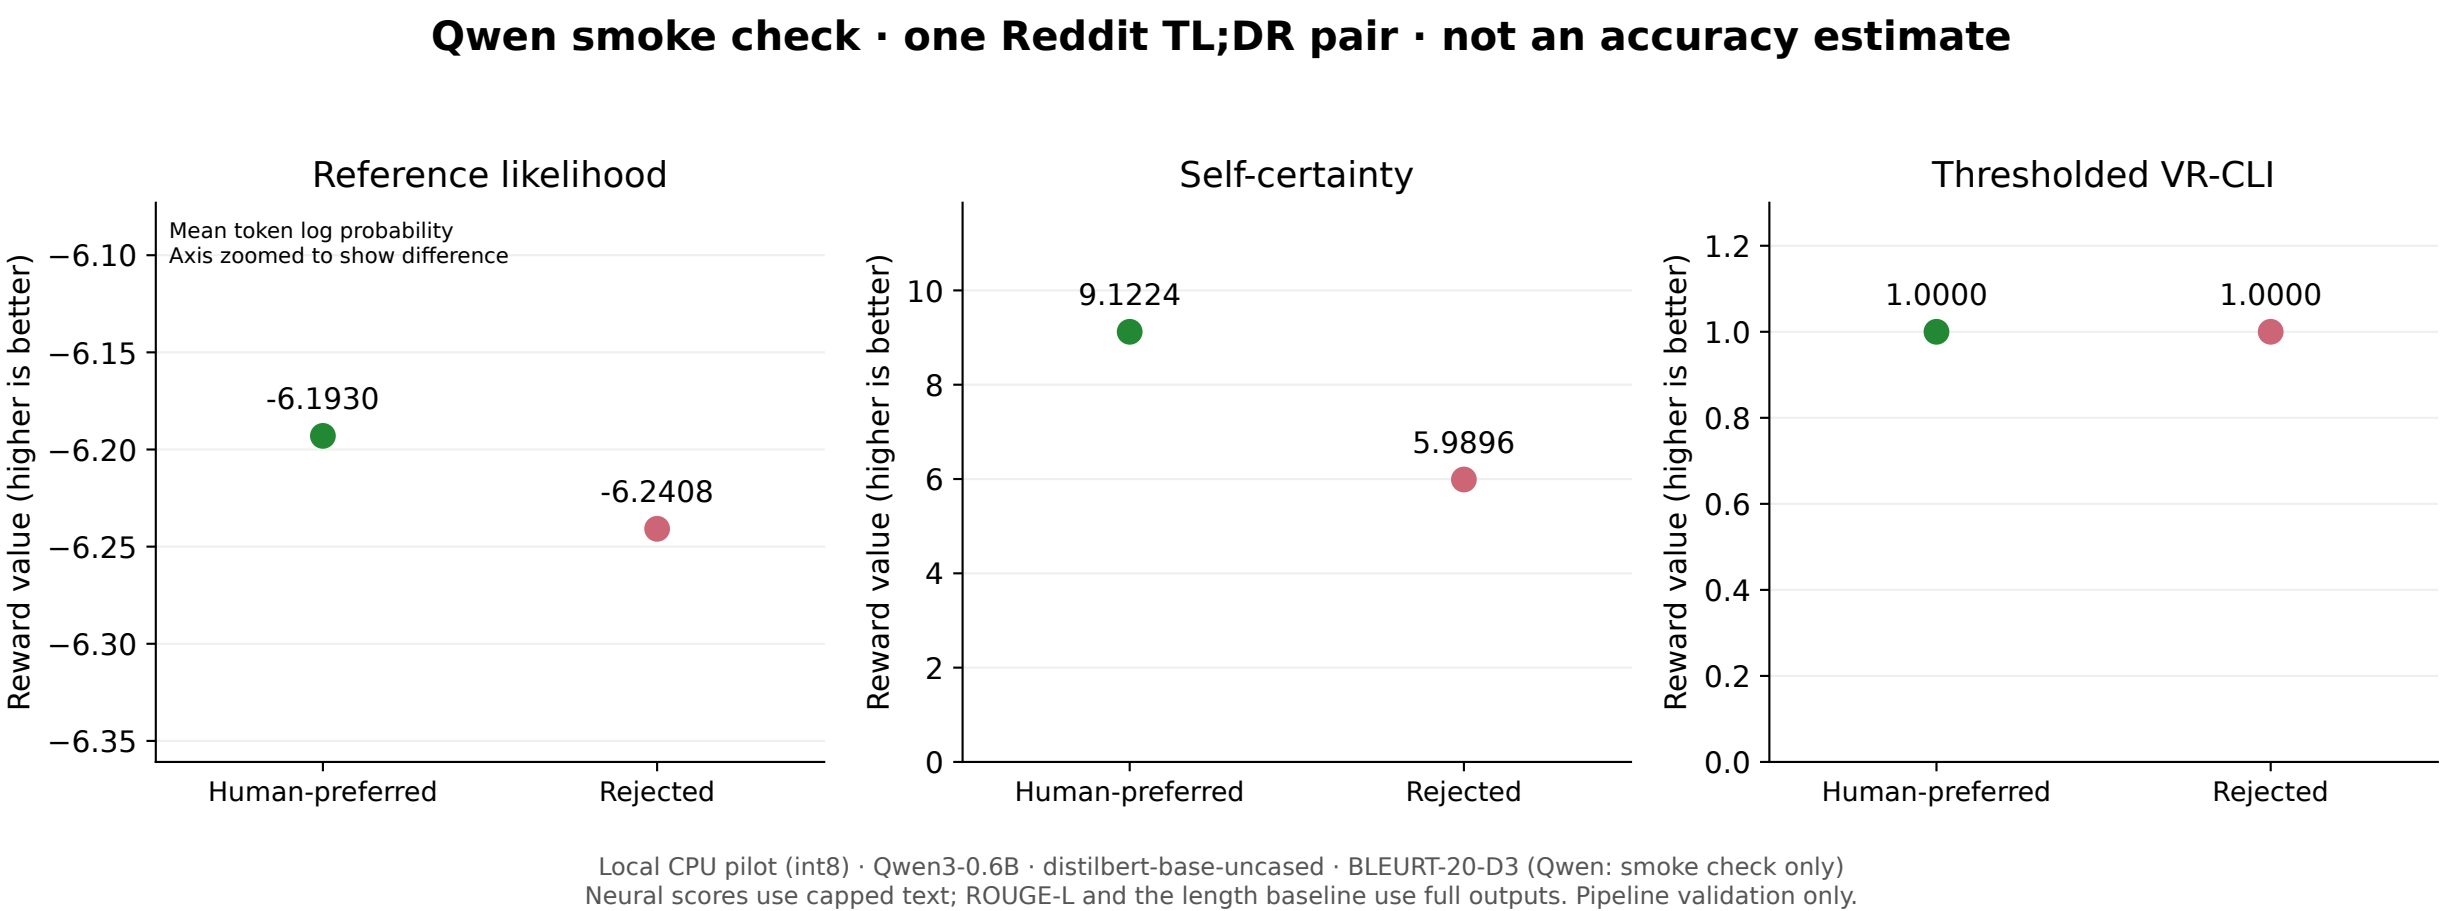
Qwen smoke check · one Reddit TL;DR pair · not an accuracy estimate
Reference likelihood
Mean token log probability
Axis zoomed to show difference
Reward value (higher is better)
−6.10
−6.15
−6.20
−6.25
−6.30
−6.35
-6.1930
-6.2408
Human-preferred
Rejected
Self-certainty
Reward value (higher is better)
10
8
6
4
2
0
9.1224
5.9896
Human-preferred
Rejected
Thresholded VR-CLI
Reward value (higher is better)
1.2
1.0
0.8
0.6
0.4
0.2
0.0
1.0000
1.0000
Human-preferred
Rejected
Local CPU pilot (int8) · Qwen3-0.6B · distilbert-base-uncased · BLEURT-20-D3 (Qwen: smoke check only)
Neural scores use capped text; ROUGE-L and the length baseline use full outputs. Pipeline validation only.